

ROYAL HOLLOWAY, UNIVERSITY OF LONDON

MSC THESIS

**Solving
the
Sliding Puzzle & Rubiks' Cube
by
Deep Reinforcement Learning**

Author:
Francois BERRIER

Supervisor:
Pr. Chris WATKINS

*A thesis submitted in fulfillment of the requirements
for the degree of MSc in Artificial Intelligence*

June 19, 2022

“The best advice I’ve ever received is ‘No one else knows what they’re doing either’”

Ricky Gervais

ROYAL HOLLOWAY, UNIVERSITY OF LONDON

Abstract

Computer Science Department

MSc in Artificial Intelligence

**Solving
the
Sliding Puzzle & Rubiks' Cube
by
Deep Reinforcement Learning**

by Francois BERRIER

The Thesis Abstract is written here (and usually kept to just this page). The page is kept centered vertically so can expand into the blank space above the title too...

Acknowledgements

blabla BofA, ST, MD, SK, CW, Vo

Contents

Abstract	v
Acknowledgements	vii
1 Objectives	1
1.1 Learning Objectives	1
1.2 Project's Objectives	2
2 Deep Reinforcement Learning	3
2.1 Reinforcement Learning	3
2.2 Deep Learning	3
2.3 Deep Reinforcement Learning	3
2.4 Graph Search & Heuristics	3
3 Games	5
3.1 Sliding Puzzle	5
3.2 Rubiks' Cube	5
4 Code	7
4.1 Code organisation	7
4.2 DRL Training	7
4.3 Solvers Comparison	7
5 Sliding Puzzle	9
5.1 2x2	9
5.2 3x3	9
5.3 4x4	9
6 Rubiks' Cube	11
6.1 2x2x2	11
6.2 3x3x3	11

List of Abbreviations

AIPnT	Artificial Intelligence Principles and Techniques
CS	Computer Science
CV	Computer Vision
DL	Deep Learning
DRL	Deep Reinforcement Learning
ML	Machine Learning
NLP	Natural Language Processing
RC	Rubiks' Cube
RL	Reinforcement Learning
RHUL	Royal Holloway, University of London
SP	Sliding Puzzle

Chapter 1

Objectives

1.1 Learning Objectives

Back when I studied Financial Mathematics, almost 2 decades ago, it was all about probability theory, stochastic calculus and asset (in particular derivatives) pricing. These skills were of course very sought after in the field of options trading, but were also often enough to get a job in algorithmic or systematic trading. By the middle of the 2010s, with the constant advances in computing power and storage, the better availability of off-the-shelves libraries and data sets, I witnessed a first revolution: the field of machine learning became more and more prominent and pretty much overshadowed other (more traditional maths) skills. More recently, a second revolution has taken not only the world of finance, but that of pretty much every science and industry, by storm: we are now in the artificial intelligence age. In 2019-2020, I decided it was time to see by myself what this was all about, and if the hype was justified. What better way to do that than embark on a proper MSc in Artificial Intelligence?

Of all the modules I have studied over the last two years of the Royal Holloway MSc in AI, I have been the most impressed by DL and NLP (itself arguably largely an application of DL) and particularly interested in AIPnT, especially our excursion in the field of graphs search (a very traditional CS topic, but which somehow I had not yet had a chance to study in much details). Even though I still believe there is a tremendous amount of malinvestment everywhere, due in good part to the inability of the average investor to distinguish between serious and scammy AI applications and startups (the same obviously goes for blockchain applications, which might warrant another MSc?), I have totally changed my mind around the potential of DL, DRL and NLP and think they are incredibly promising. I have been astonished to see by myself, through several of the courseworks we have done during the MSc, how incredibly efficient sophisticated ML, DL and DLR algorithms can be, when applied well on the right problems. Sometimes they just vastly outperform more naive and traditional approaches to the point of rendering older approaches entirely obsolete (e.g CV, NLP, game solvers, etc...).

For the project component of the MSc, I thought it would be interesting (and fun) for me to try and apply some of the DL, DRL and search techniques (from AIPnT) to a couple of single-player games, such as the sliding puzzle (of which some variations are well known under different names, e.g. the 8-puzzle and 15-puzzle) and of course the Rubik's cube. I am in particular looking to solidify my understanding of DRL by implementing and experimenting with concrete (though arguably of limited practical use) problems.

1.2 Project's Objectives

blablabla

Chapter 2

Deep Reinforcement Learning

2.1 Reinforcement Learning

blabla

2.2 Deep Learning

blabla

2.3 Deep Reinforcement Learning

blabla

2.4 Graph Search & Heuristics

See Knuth, 1998

Chapter 3

Games

3.1 Sliding Puzzle

The first puzzle I will focus on is the sliding puzzle (see Wikipedia, [2022](#))

3.2 Rubiks' Cube

blabla

Chapter 4

Code

4.1 Code organisation

blablabla

4.2 DRL Training

blabla

4.3 Solvers Comparison

blabla]

Chapter 5

Sliding Puzzle

5.1 2x2

blablabla

5.2 3x3

blabla

5.3 4x4

blabla

Chapter 6

Rubiks' Cube

6.1 2x2x2

blablabla

6.2 3x3x3

blabla

Bibliography

Books

Knuth, Donald Ervin (1998). *The art of computer programming, Volume 4, Fascile 0, Introduction to Combinatorial Algorithms and Boolean Functions, 3rd Edition*. Addison-Wesley.

Misc

Wikipedia (2022). *Sliding Puzzle*. [wikipedia](#). Accessed: 2022-06-22.